

Julia for R programmers

Douglas Bates

University of Wisconsin - Madison
<Bates@Wisc.edu>

useR!2012
June 14, 2012

1 What is Julia?

Outline

- 1 What is Julia?
- 2 The Simple Gibbs example

The Julia language

According to its developers (I have itemized what was a paragraph)

- *Julia is a high-level, high-performance dynamic programming language for technical computing, with syntax that is familiar to users of other technical computing environments.*
- *It provides a sophisticated compiler, distributed parallel execution, numerical accuracy, and an extensive mathematical function library.*
- *The library, mostly written in Julia itself, also integrates mature, best-of-breed C and Fortran libraries for linear algebra, random number generation, FFTs, and string processing.*
- *Julia programs are organized around defining functions, and overloading them for different combinations of argument types, which can also be user-defined.*

Similarities to R

- “**high-level** . . . **dynamic** programming language for **technical computing**” .
 - high-level** can work on the level of vectors, matrices, structures, etc.
 - dynamic** values have types, identifiers don't. Functions can be defined during an interactive session.
 - technical computing** these folks know about floating point arithmetic
- “organized around defining **functions**, and overloading them for different combinations of argument types”. The “overloading . . .” part means generic functions and methods.
- “syntax that is familiar to users of other technical computing environments”. Julia code looks very much like R code and/or Matlab/octave code. It is, of course, not identical but sufficiently similar to be readable.

Major differences - advantage Julia

- The ways in which R and Julia are different are too many to enumerate. These just skim the surface.
- “high-performance”, “sophisticated compiler”, “distributed parallel execution”. We will see all of these in examples. The compiler, based on LLVM, creates dynamically-loaded machine code from method instantiations - type inference is used to create very fast code.
- “types, which can be user-defined”. Julia has a sophisticated type system, simultaneously more specific and more general than classes in R. Scalar types exist (convenient for interfacing to C or Fortran compiled libraries).
- All method dispatch is S4-like multiple dispatch. All functions are generic.
- Distribution of base system and of packages is based on *git*.

Major differences - advantage R

- Named actual arguments
- Default argument values
- Handling of NA's is built-in to R at a very low level
- R is a mature language with large and active developer and user communities and an extensive infrastructure (CRAN, Bioconductor, mailing lists, The R Journal, JSS, hundreds of books describing R and its applications, ...)

The Simple Gibbs example

- To get a flavor of the language, consider an example that has been used in many language comparisons (admittedly an example not well suited to R).
- Generate a sample from the (unscaled) bivariate density

$$f(x, y) \propto x^2 \exp(-xy^2 - y^2 + 2y - 4x)$$

using direct Gibbs sampling from the conditional distributions

$$x|y \sim \Gamma(3, y^2 + 4)$$

$$y|x \sim \mathcal{N}\left(\frac{1}{1+x}, \frac{1}{2(1+x)}\right)$$

with the Γ distribution in the shape/rate formulation

- Create a function of two arguments, N and *thin*, that returns a matrix of size $N \times 2$ containing N samples after thinning by *thin*.

The R function, Rgibbs

```
Rgibbs <- function(N,thin) {  
  mat <- matrix(0,ncol=2,nrow=N)  
  x <- 0  
  y <- 0  
  for (i in 1:N) {  
    for (j in 1:thin) {  
      x <- rgamma(1,3,rate=y*y+4)  
      y <- rnorm(1,1/(x+1),1/sqrt(2*(x+1)))  
    }  
    mat[i,] <- c(x,y)  
  }  
  mat  
}
```

As stated earlier, this function will not perform well because of the inherently sequential nature of the calculation. Even the byte-compiled version is relatively slow.

JGibbs1, using Rmath.jl

```
load("extras/Rmath.jl")  
function JGibbs1(N::Int, thin::Int)  
    mat = Array{Float64, (N, 2)}  
    x = 0.  
    y = 0.  
    for i = 1:N  
        for j = 1:thin  
            x = rgamma(1,3,1/(y*y + 4))[1]  
            y = rnorm(1, 1/(x+1),1/sqrt(2(x + 1)))[1]  
        end  
        mat[i,:] = [x,y]  
    end  
    mat  
end
```

The arguments to `rgamma` here are shape and scale (not rate).

JGibbs2, direct calls to libRmath

```
function JGibbs2(N::Int, thin::Int)
    mat = Array{Float64, (N, 2)}
    x = 0.
    y = 0.
    for i = 1:N
        for j = 1:thin
            x = ccall(dlsym(_jl_libRmath, :rgamma),
                      Float64, (Float64, Float64),
                      3., 1/(y*y + 4))
            y = ccall(dlsym(_jl_libRmath, :rnorm),
                      Float64, (Float64, Float64),
                      1/(x+1), 1/sqrt(2*(x + 1)))
        end
        mat[i, :] = [x, y]
    end
    mat
end
```

JGibbs3, the Julia randn and randg samplers

```
function JGibbs3(N::Int , thin::Int)
    mat = Array{Float64} (N, 2)
    x    = 0.
    y    = 0.
    for i = 1:N
        for j = 1:thin
            x = randg(3) / (y*y + 4)
            y = 1/(x + 1) + randn()/sqrt(2(x + 1))
        end
        mat[i ,:] = [x,y]
    end
    mat
end
```

Distributed versions, dJGibbs3a and dJGibbs3b

- Julia allows for distributed parallel execution by specifying the number of processes at start-up.
- One appealing abstraction for parallel execution is a “distributed array” where each process works on a part of an array.
- The dJGibbs3a function leaves the result as a distributed array, suitable for further distributed processing. dJGibbs3b returns the result as an ordinary array. The results are from multiple chains not a single long chain as in the other functions.

```
function dJGibbs3a(N::Int, thin::Int)
    darray((T,d,da)→JGibbs3(d[1],thin),
           Float64, (N, 2), 1)
end

function dJGibbs3b(N::Int, thin::Int)
    convert(Array{Float64,2}, dJGibbs3a(N, thin))
end
```

Timings

- Detailed timings (and the code if you want to try yourself) are at <https://gist.github.com/2656226>.
- Roughly the results are:
 - ▶ JGibbs1 is within a factor of 2 of RcppGibbs.
 - ▶ JGibbs2 is nearly the same speed as RcppGibbs.
 - ▶ JGibbs3 is faster than RcppGibbs (which uses the R samplers) and GSLGibbs (using the GSL samplers). Differences are attributable to different samplers.
 - ▶ On a 4-core processor using 4 processes (and no conflicting jobs), dJGibbs3a is nearly 4x faster than JGibbs3. dJGibbs3b is about 3x faster than JGibbs3, due to the communication overhead of converting from distributed to non-distributed.

Recall the comment about “library written in Julia itself”

```
# Generating gamma variables – Marsaglia and Tsang
function randg(a::Real)
    d = a - 1.0/3.0
    c = 1.0 / sqrt(9*d)
    while(true)
        v = 0.
        while (v <= 0.0)
            x = randn()
            v = 1.0 + c*x
        end
        v = v*v*v
        U = rand()
        x2 = x*x
        if U < 1.0 - 0.331*x2*x2; return d*v; end
        if log(U) < 0.5*x2 + d*(1.0 - v + log(v))
            return d*v
        end
    end
end
```